



COL333/671: Introduction to AI

Semester I, 2022-23

Adversarial Search

Rohan Paul

Outline



- Last Class
 - Constraint Satisfaction
- This Class
 - Adversarial Search
- Reference Material
 - AIMA Ch. 5 (Sec: 5.1-5.5)

Acknowledgement

These slides are intended for teaching purposes only. Some material has been used/adapted from web sources and from slides by Doina Precup, Dorsa Sadigh, Percy Liang, Mausam, Dan Klein, Anca Dragan, Nicholas Roy and others.

Game Playing and AI

- **Games: challenging decision-making problems**
 - Incorporate the state of the other agent in your decision-making. Leads to a vast number of possibilities.
 - Long duration of play. Win at the end.
 - Time limits: Do not have time to compute optimal solutions.



Games: Characteristics

- Axes:
 - Players: one, two or more.
 - Actions (moves): deterministic or stochastic
 - States: fully known or not.
- Zero-Sum Games
 - Adversarial: agents have opposite utilities (values on outcomes)

- **Core: contingency problem**

- The opponent's move is **not** known ahead of time. A player must respond with a move for **every possible** opponent reply.

- **Output**

- Calculate a **strategy (policy)** which recommends a move from each state.

Playing Tic-Tac-Toe: *Essentially a search problem!*

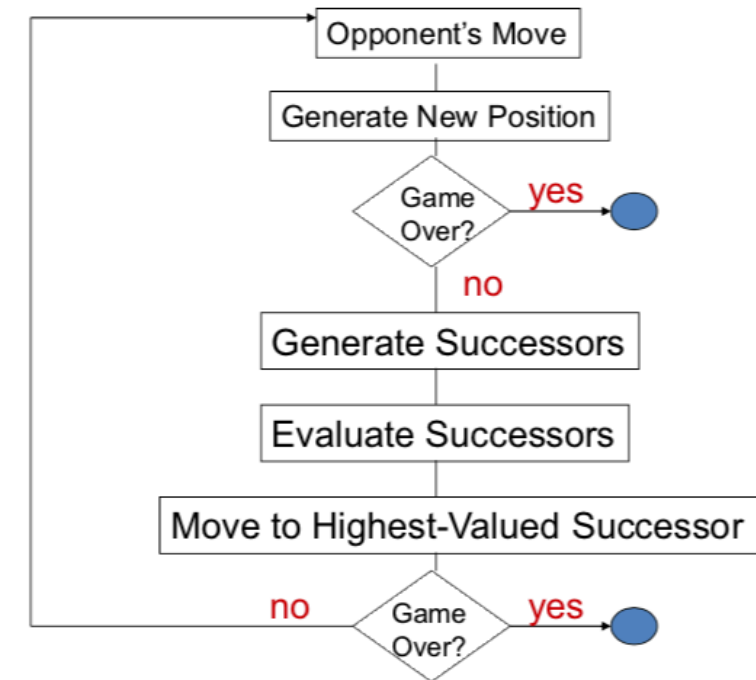
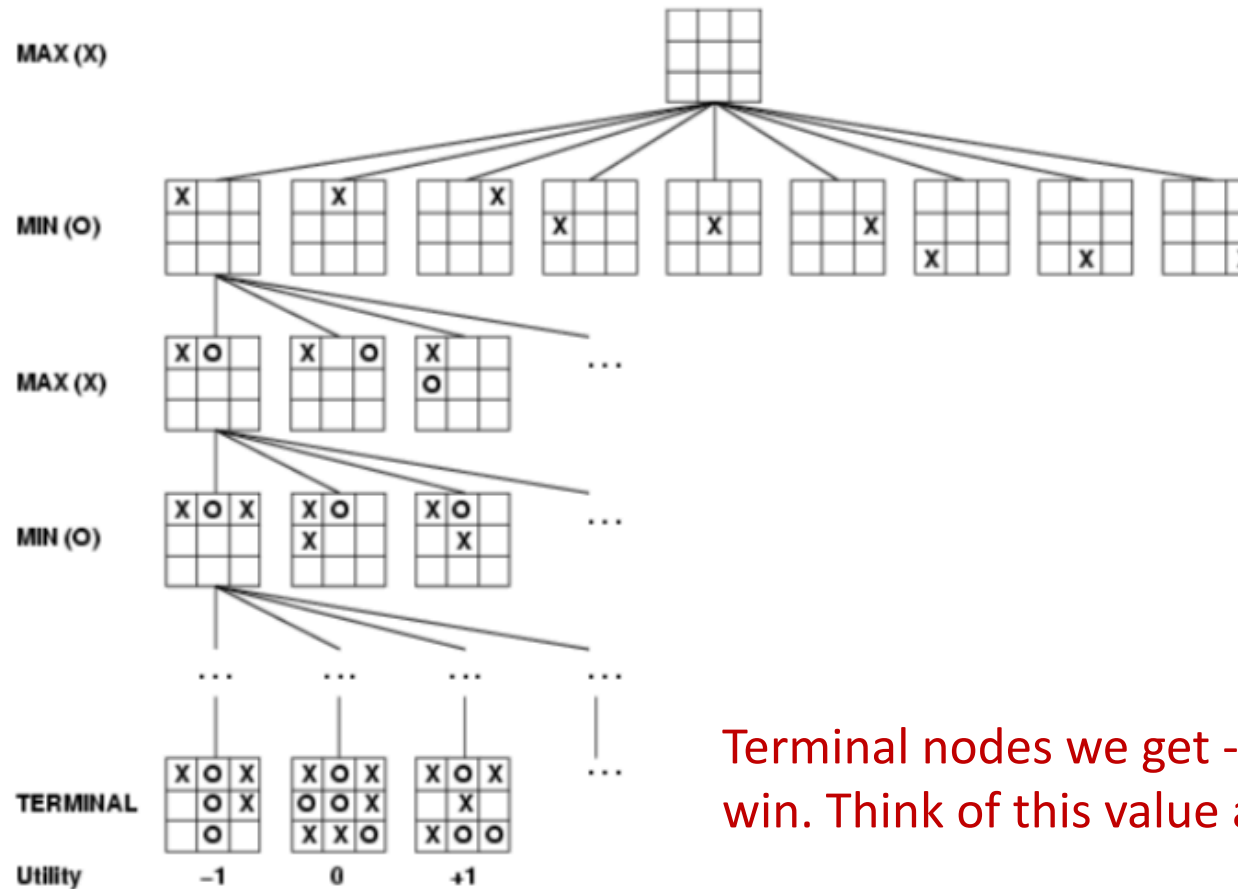
computer's
turn

opponent's
turn

computer's
turn


opponent's
turn

leaf nodes
are evaluated




Terminal nodes we get -1, 0 or 1 for loss, tie or win. Think of this value as a "utility" of a state.

Single-Agent Trees




Computing “utility” of states to decide actions

Value of a state:
The best achievable
outcome (utility)
from that state



Game Trees: Presence of an Adversary



The adversary's actions are not in our control. Plan as a contingency considering all possible actions taken by the adversary.


Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:


$$V(s) = \text{known}$$

Adversarial Search (Minimax)


- Consider a deterministic, zero-sum game
 - Tic-tac-toe, chess etc.
 - One player maximizes result and the other minimizes result.
- Minimax Search
 - Search the game tree for best moves.
 - Select optimal actions that move to a position with the highest minimax value.
 - What is the minimax value?
 - It is the best achievable utility against the optimal (rational) adversary.
 - Best achievable payoff against the best play by the adversary.

Minimax Algorithm

- Ply and Move
 - Move: when action taken by both players.
 - Ply: is a half move.
- Backed-up value
 - of a MAX-position: the value of the largest successor
 - of a MIN-position: the value of its smallest successor.
- Minimax algorithm
 - Search down the tree till the terminal nodes.
 - At the bottom level apply the utility function.
 - Back up the values up to the root along the search path (compute as per min and max nodes)
 - The root node selects the action.



Minimax Example



Minimax Implementation


def max-value(state):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

 return v



def min-value(state):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

 return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Implementation

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

```
def min-value(state):
```

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}))$

return v

Useful, when there are multiple adversaries.

Minimax Properties

- Completeness

- Yes

- Complexity

- Time: $O(b^m)$
 - Space: $O(bm)$

- Requires growing the tree till the terminal nodes.
 - Not feasible in practice for a game like Chess.

- Chess:

- branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^m \approx 35^{100} \approx 10^{154}$

- The Universe:



- number of atoms $\approx 10^{78}$
 - age $\approx 10^{18}$ seconds
 - 10^8 moves/sec $\times 10^{78} \times 10^{18} = 10^{104}$

Minimax Properties

- Optimal


- If the adversary is playing optimally (i.e., giving us the min value)
 - Yes
- If the adversary is not playing optimally (i.e., not giving us the min value)
 - No. Why? It does not exploit the opponent's weakness against a suboptimal opponent).

You: Cricle. Opponent: Cross



If min returns 9? Or 100?


Necessary to examine all values in the tree?



Alpha-Beta Pruning: General Idea

- **General Configuration (MIN version)**


- Consider computing the MIN-VALUE at some node n , examining n 's children
- n 's estimate of the childrens' min is reducing.
- Who can use n 's value to make a choice? MAX
- Let a be the best value that MAX can get at any choice point along the current path from the root
- If the value at n becomes worse than a , MAX will not pick this option, so we can stop considering n 's other children (any further exploration of children will only reduce the value further)




Alpha-Beta Pruning: General Idea

- **General Configuration (MAX version)**


- Consider computing the MAX-VALUE at some node n , examining n 's children
- n 's estimate of the childrens' max is increasing.
- Who can use n 's value to make a choice? MIN
- Let b be the lowest (best) value that MIN can get at any choice point along the current path from the root
- If the value at n becomes higher than b , MIN will not pick this option, so we can stop considering n 's other children (any further exploration of children will only increase the value further)



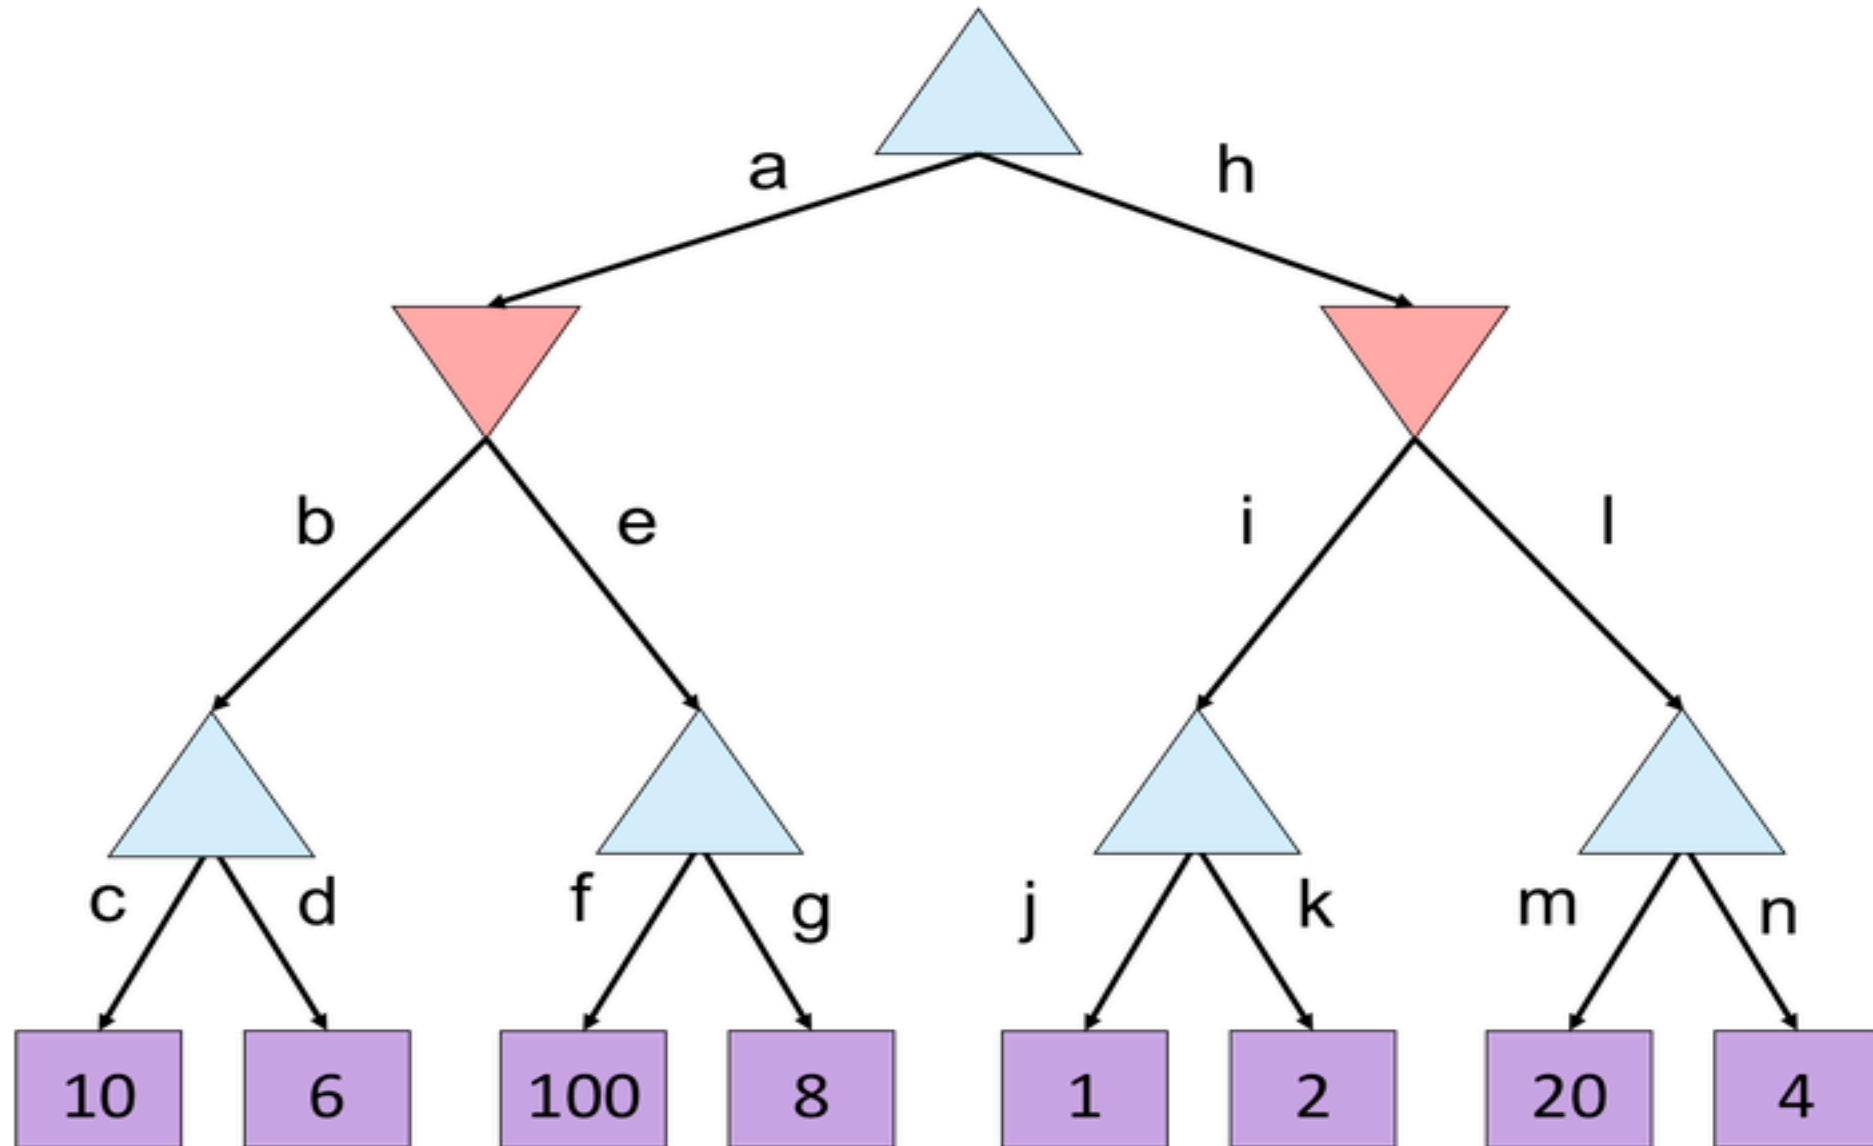
Pruning: Example




Pruning: Example



Pruning: Example



Pruning: Example



Alpha-Beta Implementation

α : MAX's best option on path to root
 β : MIN's best option on path to root

def max-value(state, α , β):

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \geq \beta$ return v

$\alpha = \max(\alpha, v)$

return v

def min-value(state, α , β):

initialize $v = +\infty$

for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

if $v \leq \alpha$ return v


$\beta = \min(\beta, v)$

return v


Alpha-Beta Pruning - Properties

1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.

Alpha-Beta Pruning – Order of nodes matters




Alpha-Beta Pruning – Order of nodes matters




Alpha-Beta Pruning - Properties

1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.
3. The alpha-beta search cuts the largest amount off the tree when we examine the **best move first**
 - However, best moves are typically **not** known. Need to make estimates.

Alpha-Beta Pruning – Order of nodes matters



If the nodes were indeed encountered as “worst moves first” – then no pruning is possible



If the nodes were encountered as “best moves first” – then pruning is possible

Note: In reality, we don’t know the ordering.

Alpha-Beta Pruning - Properties

1. Pruning has **no effect** on the minimax value at the root.
 - Pruning does not affect the final action selected at the root.
2. A form of **meta-reasoning** (computing what to compute)
 - Eliminates nodes that are irrelevant for the final decision.
3. The alpha-beta search cuts the largest amount off the tree when we examine the **best move first**
 - Problem: However, best moves are typically **not** known.
 - Solution: Perform iterative deepening search and evaluate the states.
4. Time Complexity
 - **Best ordering** - $O(b^{m/2})$. Can double the search depth for the same resources. Effective branching factor becomes $b^{1/2}$ instead of b .
 - On average – $O(b^{3m/4})$ if we expect to find the min or max after $b/2$ expansions.

Minimax for Chess

- Chess:
 - branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^m \approx 35^{100} \approx 10^{154}$
- The Universe:
 - number of atoms $\approx 10^{78}$
 - age $\approx 10^{18}$ seconds
 - 10^8 moves/sec $\times 10^{78} \times 10^{18} = 10^{104}$

Alpha-Beta for Chess


- Chess:
 - branching factor $b \approx 35$
 - game length $m \approx 100$
 - search space $b^{m/2} \approx 35^{50} \approx 10^{77}$

Cutting-off Search

- Problem (Resource constraint):
 - Minimax search: full tree till the terminal nodes.
 - Alpha-beta prunes the tree but still searches till the terminal nodes.
 - We can't search till the terminal nodes.
- Solution:
 - Depth-limited Search (H-Minimax)
 - Search only to a limited depth (cutoff) in the tree
 - **Replace the terminal utilities with an evaluation function for non-terminal positions.**


H-MINIMAX(s, d) =

$$\begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN.} \end{cases}$$



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search.
- Estimate the chances of winning.



- Ideal function: returns the actual **minimax** value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$


- e.g. $f_i(s)$ = (number of pieces of type i), each weight w_i etc.

Evaluation Functions


- Evaluation functions take a state and output an estimate of the true minimax value of that node.
 - Typically, “better” states will be assigned higher values by a good evaluation function in comparison to “worse” states. Evaluation functions serve a similar purpose as heuristics in classical search.
- Depth-limited search applies evaluation function at the maximum solvable depth
 - Gives them mock terminal utilities by the evaluation function.
- Evaluation functions require features (some aspect of the current state).
 - Functions may or may not be linear. Require considerable thought and experimentation for designing.
- The better the evaluation function is, the closer the agent will come to behaving optimally.
 - Going deeper into the tree before using the evaluation function also tends to give better results. Reduces the compromise of optimality.

Determining “good” node orderings

- The ordering of nodes helps alpha-beta pruning.
 - Worst ordering $O(b^m)$. Best ordering $O(b^{m/2})$.
- How to find good orderings
 - Problem: we only know them when we evaluate the nodes.




- One approach – iterative deepening to determine evaluations for nodes
 - What if we can do iterative deepening to a certain depth. Use the evaluation function at the set depth and then compute the values for the nodes in the tree that is generated.
 - Next time, use the evaluations of the previous search to order the nodes. Use them for pruning.
 - Use evaluations of the previous search for order.



Game of Chance: Expectimax

- When the result of an action is not exactly known. Need a notion of uncertainty or chance in action selection.
- Explicit randomness in the opponent's action selection
 - Unpredictable opponents: the ghosts move randomly in Pacman.
 - Rolling dice by a player in a game.
- Pessimistic assumption is not valid for the adversary
 - The adversary may not be that bad. May not provide the worst value. Optimal response may not be guaranteed.



Expectimax:


At chance nodes the outcome is uncertain. Calculate the **expected utilities**: weighted average (expectation) of children

Expectimax Search


$$\forall \text{ agent-controlled states, } V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

$$\forall \text{ chance states, } V(s) = \sum_{s' \in \text{successors}(s)} p(s'|s) V(s')$$

$$\forall \text{ terminal states, } V(s) = \text{known}$$





$$v = (1/2) (8) + (1/3) (24) + (1/6) (-12) = 10$$



Mixed-type layers in a game tree are also possible. More than two agents.

Expectimax Search



Can we perform pruning?

Expectimax Search

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is EXP: return exp-value(state)

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{value}(\text{successor}))$

return v

```
def exp-value(state):
```

initialize $v = 0$

for each successor of state:


$p = \text{probability}(\text{successor})$

$v += p * \text{value}(\text{successor})$

return v


Depth-Limited Expectimax

- Depth-limit can also be applied in Expectimax search.
- Use heuristics to estimate the values at the depth limit.



Example: Game of Go

- The game of Go originated in China more than 2000 years ago.
- Usually played on 19x19, also 13x13 or 9x9 board
- Black and white place down stones alternately.
- Surrounded stones are captured and removed.
- The player with more territory wins the game.
- Complex strategy for capturing and creating a territory.
- Grand challenge in AI game playing because of its complexity.



Example: Game of Go

Significantly higher branching factor compared to Chess.

- Alpha-beta pruning/minimax does not scale. Not easy to evaluate all the action outcomes.

Design of a heuristic function is difficult

- Most positions are in a flux till the end game. Value not a strong indicator of winning.

Alternate approach, Monte-Carlo Tree Search.

Popularized by Alpha Go

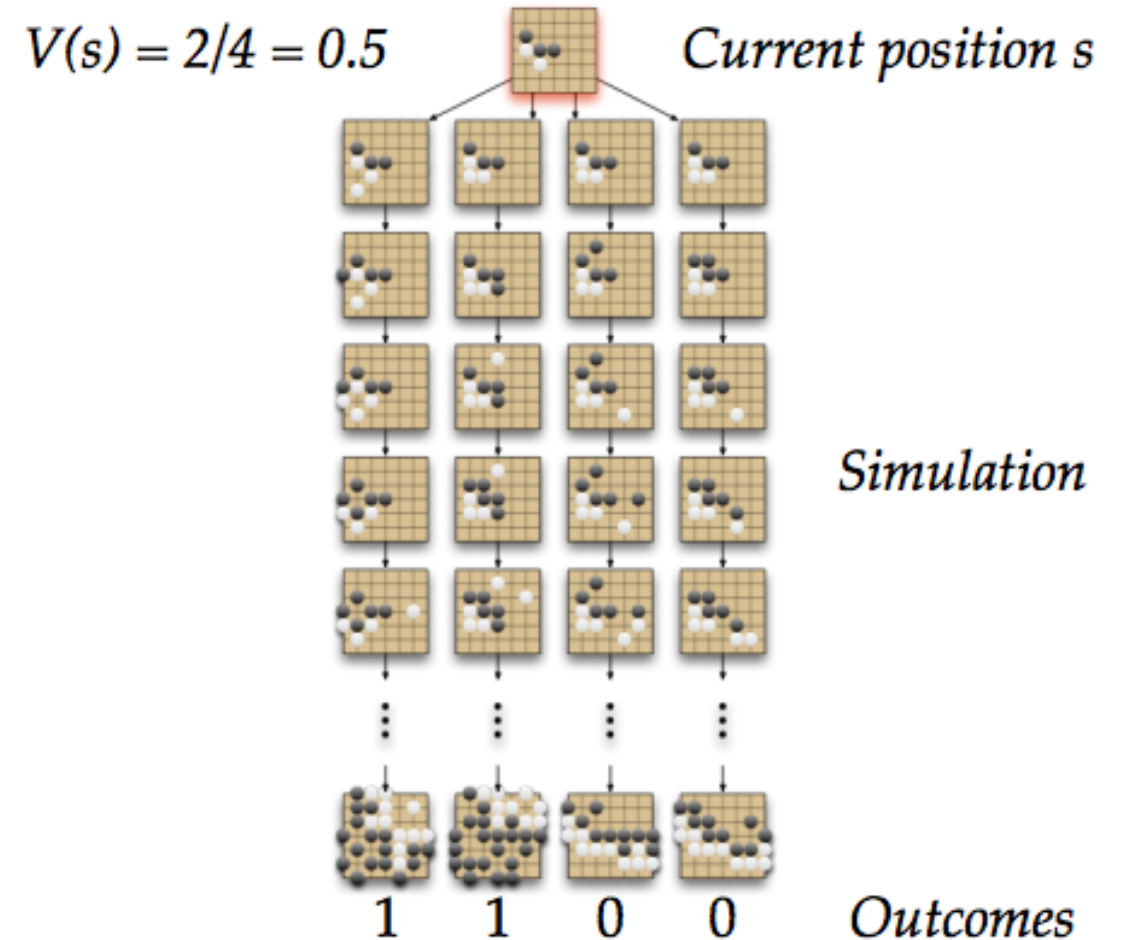
<https://www.deepmind.com/research/highlighted-research/alphago>

	Chess	Go
Size of board	8 x 8	19 x 19
Average no. of moves per game	100	300
Avg branching factor per turn	35	235
Additional complexity		Players can pass

Monte Carlo Tree Search (MCTS)

1. Simulations/Rollouts



- Evaluation of a state $V(s)$ using roll outs or simulating what will happen from this state onwards.
 - From state s play many times using a policy (e.g., random) and count wins and losses.
- For games in which the only outcomes are a win or a loss,
 - The “win percentage” approximates the “average utility”.




Monte Carlo Tree Search (MCTS)

2. Selective Search

- May not evaluate all states.
 - Be selective with evaluations on more promising actions/states.
- Explore parts of the tree (without an explicit depth for exploration) that will
 - Improve the decision at the root (improve the estimation of the value function)
 - Grow the tree of states as needed to improve the value estimates of a state.



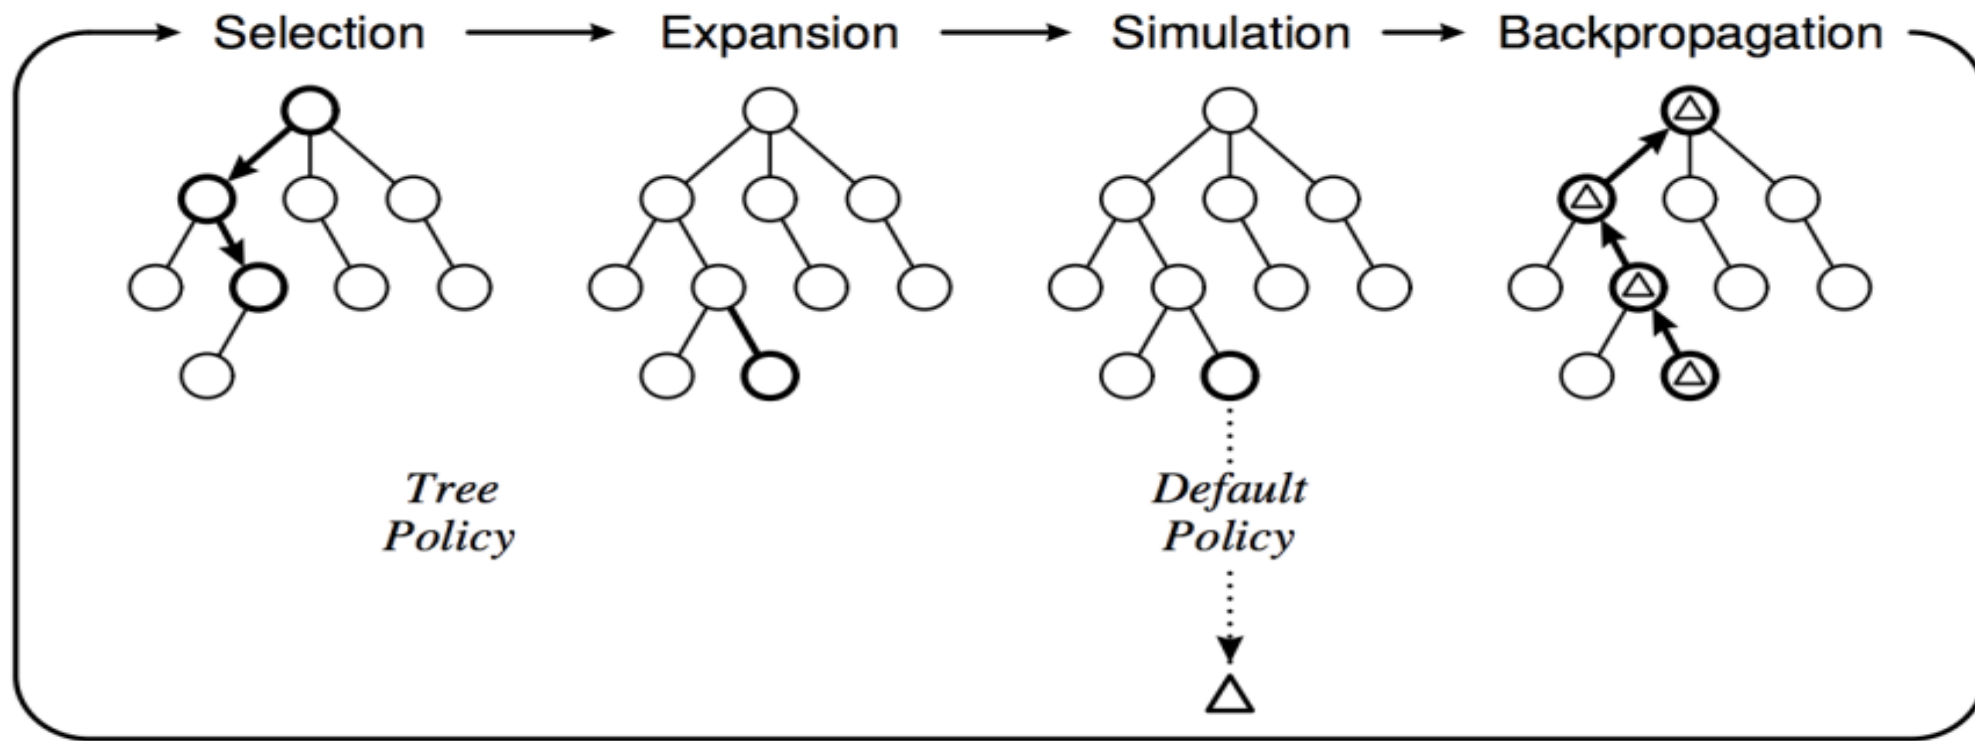


Selection

- Start from the root and select a move (via a selection/tree policy).
- Used for nodes we *have* seen before

Expansion

- When we reach the frontier, grow the search tree by generating a new child node of the node selected from the frontier.



Selection

- Start from the root and select a move (via a selection/tree policy).
- Used for nodes we *have* seen before

Expansion

- When we reach the frontier, grow the search tree by generating a new child node of the node selected from the frontier.


Simulation

- Perform playout from the newly generated child node.
- Select moves for both players according to a playout policy (also called default policy) such as random action selection.
- Do not record the nodes in the tree.


Backpropagation

- After reaching a terminal node
- Update value and visits for states expanded in selection and expansion


Example



(a) Selection



(b) Expansion and simulation



(c) Backpropagation


MCTS Procedure

```
function MONTE-CARLO-TREE-SEARCH(state) returns an action  
  tree  $\leftarrow$  NODE(state)  
  while IS-TIME-REMAINING() do  
    leaf  $\leftarrow$  SELECT(tree)  
    child  $\leftarrow$  EXPAND(leaf)  
    result  $\leftarrow$  SIMULATE(child)  
    BACK-PROPAGATE(result, child)  
  return the move in ACTIONS(state) whose node has highest number of playouts
```

Exploration vs. Exploitation

Selection Strategy

- How to select moves/actions in the tree?
- Bias the moves towards those providing higher value.
- But we may not know about the value of certain states or may be very uncertain about them. Hence, sometimes we should explore too.
- Fundamental trade-off between exploration and exploitation.




How to select the moves balancing exploration and exploitation.

Upper Confidence Bound applied to Trees



$$UCB1(n) = \frac{U(n)}{N(n)} + C \times \sqrt{\frac{\log N(\text{PARENT}(n))}{N(n)}}$$

- $N(n)$ = number of rollouts from node n
- $U(n)$ = total utility of rollouts (e.g., # wins) for $\text{Player}(\text{Parent}(n))$
- C is the tunable parameter.
- The first term is the exploitation term: the average utility of node n .
- The second term is the exploration term: how uncertain we are about the node's utility.
 - The denominator is the number of visits to the states, so states visited less often are preferred.
 - The numerator is the log of the number of times the parent is explored.
 - If we are selecting n for some non-zero percentage of times then the exploration term goes to zero as the counts increase.
- We will revisit this concept in the discussion on Reinforcement Learning later.

AlphaGo Team (Silver et al., 2016)



adapted from Sylvain Gelly & David Silver, Test of Time Award ICML 2017




(Silver et al., 2016)

Alpha Go combined learning with MCTS (used a NN to predict values/utilities of states). Employed self play etc.

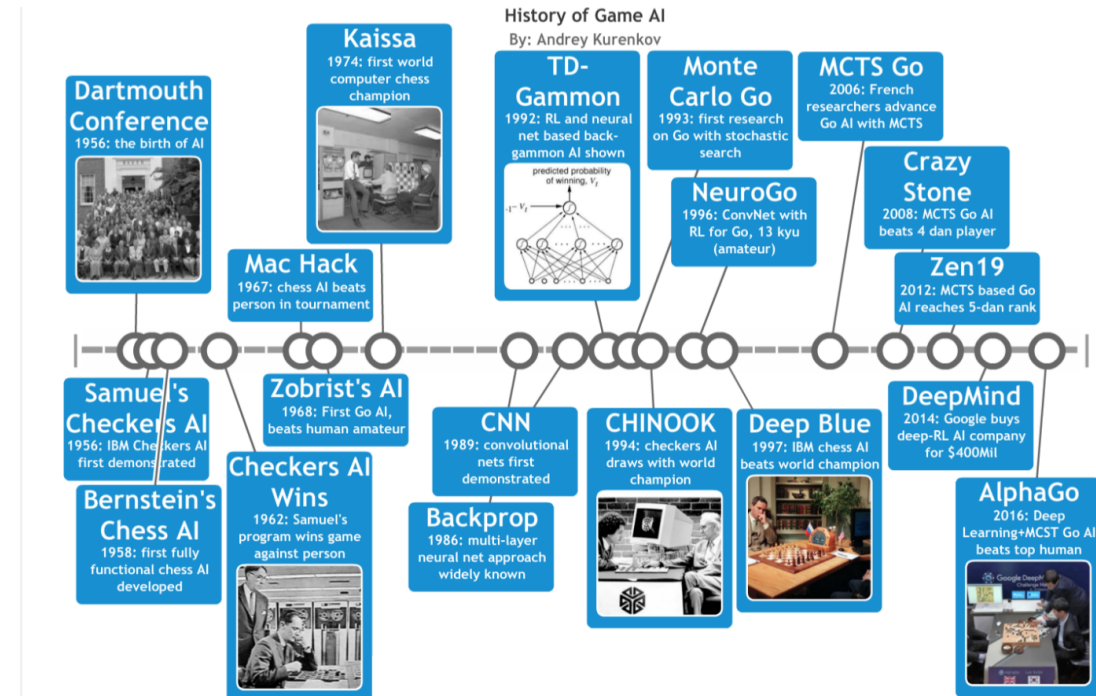
Multiple players and other games

- Not all games are zero sum.
 - Loss for one agent may not be win for the other agent.
 - Different agents may have different tasks in the game that don't directly involve strictly competing against each other.
- Multi-agent utilities.
 - Generalization of minimax.
 - Each player maximizes its own utility at each node they control and ignore the utilities of the other agents.
- General games with multi-agent utilities
 - Can invoke cooperation
 - The utility selected at the root tends to yield a reasonable utility for all participating agents.



Game Playing AI: Wrap up

- Game playing domains
 - Very large amount of contingency reasoning.
- Exact decision making is nearly impossible.
 - Approximate evaluation functions etc.
 - Force efficient use of computation (alpha-beta pruning.)
- An important test bed for AI algorithms.
 - We play games intuitively, used to reasoning.
 - Easy to compare human and computer performance.
- Game playing has produced important research ideas
 - Reinforcement learning (checkers)
 - Iterative deepening (chess)
 - Monte Carlo tree search (chess, Go)
 - Solution methods for partial-information games in economics (poker)



*"Games are to AI as grand prix is to automobile design"
Games viewed as an indicator of intelligence.*